

System Programming Lab #1

2019-03-12

sp-tas

Case Study: Library Interpositioning

- **Library interpositioning : powerful linking technique that allows programmers to intercept calls to arbitrary functions**
- **Interpositioning can occur at:**
 - Compile time: When the source code is compiled
 - Link time: When the relocatable object files are statically linked to form an executable object file
 - Load/run time: When an executable object file is loaded into memory, dynamically linked, and then executed.

Some Interpositioning Applications

■ Security

- Confinement (sandboxing)
 - Interpose calls to libc functions.
- Behind the scenes encryption
 - Automatically encrypt otherwise unencrypted network connections.

■ Monitoring and Profiling

- Count number of calls to functions
- Characterize call sites and arguments to functions
- Malloc tracing
 - Detecting memory leaks
 - **Generating address traces**

Example program

```
#include <stdio.h>
#include <stdlib.h>
#include <malloc.h>

int main()
{
    free(malloc(10));
    printf("hello, world\n");
    exit(0);
}
```

hello.c

- **Goal: trace the addresses and sizes of the allocated and freed blocks, without modifying the source code.**
- **Three solutions: interpose on the `lib malloc` and `free` functions at compile time, link time, and load/run time.**

Compile-time Interpositioning

```
#ifdef COMPILETIME
/* Compile-time interposition of malloc and free using C
 * preprocessor. A local malloc.h file defines malloc (free)
 * as wrappers mymalloc (myfree) respectively.
 */

#include <stdio.h>
#include <malloc.h>

/*
 * mymalloc - malloc wrapper function
 */
void *mymalloc(size_t size, char *file, int line)
{
    void *ptr = malloc(size);
    printf("%s:%d: malloc(%d)=%p\n", file, line, (int)size,
ptr);
    return ptr;
}
```

mymalloc.c

Compile-time Interpositioning

```
#define malloc(size) mymalloc(size, __FILE__, __LINE__ )
#define free(ptr) myfree(ptr, __FILE__, __LINE__ )

void *mymalloc(size_t size, char *file, int line);
void myfree(void *ptr, char *file, int line);
```

`malloc.h`

```
linux> make helloc
gcc -O2 -Wall -DCOMPILETIME -c mymalloc.c
gcc -O2 -Wall -I. -o helloc hello.c mymalloc.o
linux> make runc
./helloc
hello.c:7: malloc(10)=0x501010
hello.c:7: free(0x501010)
hello, world
```

Link-time Interpositioning

```
#ifdef LINKTIME
/* Link-time interposition of malloc and free using the
static linker's (ld) "--wrap symbol" flag. */

#include <stdio.h>

void *__real_malloc(size_t size);
void __real_free(void *ptr);

/*
 * __wrap_malloc - malloc wrapper function
 */
void *__wrap_malloc(size_t size)
{
    void *ptr = __real_malloc(size);
    printf("malloc(%d) = %p\n", (int)size, ptr);
    return ptr;
}
```

mymalloc.c

Link-time Interpositioning

```
linux> make hello1
gcc -O2 -Wall -DLINKTIME -c mymalloc.c
gcc -O2 -Wall -Wl,--wrap,malloc -Wl,--wrap,free \
-o hello1 hello.c mymalloc.o
linux> make run1
./hello1
malloc(10) = 0x501010
free(0x501010)
hello, world
```

- The “-Wl” flag passes argument to linker
- Telling linker “--wrap,malloc” tells it to resolve references in a special way:
 - Refs to `malloc` should be resolved as `__wrap_malloc`
 - Refs to `__real_malloc` should be resolved as `malloc`


```
#ifdef RUNTIME
/* Run-time interposition of malloc and free based on
 * dynamic linker's (ld-linux.so) LD_PRELOAD mechanism */
#define _GNU_SOURCE
#include <stdio.h>
#include <stdlib.h>
#include <dlfcn.h>

void *malloc(size_t size)
{
    void *(*mallocp)(size_t size);
    char *error;
    void *ptr;

    /* get address of libc malloc */
    if (!mallocp) {
        mallocp = dlsym(RTLD_NEXT, "malloc");
        if ((error = dlerror()) != NULL) {
            fputs(error, stderr);
            exit(1);
        }
    }
    ptr = mallocp(size);
    fprintf(stderr, "malloc(%d) = %p\n", (int)size, ptr);
    return ptr;
}
```

Load/Run-time Interpositioning

mymalloc.c

Load/Run-time Interpositioning

```
linux> make hellor
gcc -O2 -Wall -DRUNTIME -shared -fPIC -o mymalloc.so mymalloc.c
gcc -O2 -Wall -o hellor hello.c
linux> make runr
(LD_PRELOAD="/usr/lib/x86_64-linux-gnu/libdl.so ./mymalloc.so"
./hellor)
malloc(10) = 0x559a34eca260
free(0x559a34eca260)
hello, world
```

- The `LD_PRELOAD` environment variable tells the dynamic linker to resolve unresolved refs (e.g., to `malloc`) by looking in `libdl.so` and `mymalloc.so` first.
 - `libdl.so` necessary to resolve references to the `dlopen` functions.
- `-shared`: telling linker to make output as a shared objective (`.so`)
- `-fPIC`: Position-Independent Code

Interpositioning Recap

■ Compile Time

- Apparent calls to malloc/free get macro-expanded into calls to mymalloc/myfree

■ Link Time

- Use linker trick to have special name resolutions
 - malloc → __wrap_malloc
 - __real_malloc → malloc

■ Compile Time

- Implement custom version of malloc/free that use dynamic linking to load library malloc/free under different names

Dynamic Memory Management

`void *malloc(size_t size)`

`malloc` allocates `size` bytes of memory on the process' heap and returns a pointer to it that can subsequently be used by the process to hold up to `size` bytes. The contents of the memory are undefined.

`void *calloc(size_t nmemb, size_t size)`

`calloc` allocates `nmemb*size` bytes of memory on the process' heap and returns a pointer to it that can subsequently be used by the process to hold up to `size` bytes. The contents of the memory are set to zero.

`void *realloc(void *ptr, size_t size)`

`realloc` changes the size of the memory block pointed to by `ptr` to `size` bytes. The contents are copied up to `min(size, old size)`, the rest is undefined.

`void free(void *ptr)`

`free` explicitly frees a previously allocated block of memory.

Dynamic Memory Management

```
#include <stdlib.h>

void main(void) {
    void *p;
    char *str;
    int *A;

    // allocated 1024 bytes of memory
    p = malloc(1024);

    // allocated an integer array with 500 integer
    A = (int*)calloc(500, sizeof(int));

    // allocate a string with 16 characters...
    str = (char*)malloc(16*sizeof(char));

    // ...then resize that string to hold 512 characters
    str = (char*)realloc(str, 512*sizeof(char));

    // finally, free all allocated memory
    free(p);
    free(A);
    free(str);
}
```

example1.c

Shared library loading interfaces

- `#include <dlfcn.h>`
- `void *dlopen(const char *filename, int flags)`
 - `dlopen` loads the dynamic shared object (shared library) file named by the null-terminated string `filename` and returns an opaque “handle” for the loaded object.
 - **RTLD_LAZY** – perform lazy binding
 - **RTLD_NOW** – all undefined symbols in the shared object are resolved before `dlopen` returns
- `void *dlsym(void *handle, const char *symbol)`
 - `dlsym` takes a “handle” of a dynamic loaded shared object returned by `dlopen` and returns the address where that symbol is loaded into memory
 - **RTLD_DEFAULT** – find the first occurrence of the desired symbol
 - **RTLD_NEXT** – find the next occurrence of the desired symbol in the search order after the current object
- `int dlclose(void *handle)`
 - `dlclose` decrements the reference count on the dynamically loaded shared object referred to by `handle`

Lab Assignment #1 – Linker Lab

- Download skeleton code from eTL
- Hand In
 - Upload your files **eTL**
 - A tarball of your implementation (20/20/20/+10 pts for each part)
 - A report (10 pts)
- PLEASE, **READ** the Hand-out!!!
- Assigned: Mar. 12
- Deadline: Mar. 26, 11:59:59 PM
- Lab #2 (3/19) will be Q&A session

(Part 1) Tracing dynamic memory allocation

test1.c

```
#include <stdlib.h>

void main(void) {
    void *a;

    a = malloc(1024);
    a = malloc(32);
    free(malloc(1));
    free(a);
}
```

output

```
stuXXX@spN ~/linklab/part1 $ make run test1
[0001] Memory tracer started.
[0002]          (nil) : malloc( 1024 ) = 0xb87010
[0003]          (nil) : malloc( 32 ) = 0xb87420
[0004]          (nil) : malloc( 1 ) = 0xb87450
[0005]          (nil) : free( 0xb87450 )
[0006]          (nil) : free( 0xb87420 )
[0007]
[0008] Statistics
[0009]   allocated_total      1057
[0010]   allocated_avg        352
[0011]   freed_total         0
[0012]
[0013] Memory tracer stopped.
stuXXX@spN ~/linklab/part1 $
```


(Part 2) Tracing unfreed memory

test1.c

```
#include <stdlib.h>
```

```
void main(void) {  
    void *a;
```

output

```
    a = malloc(1024);  
    a = malloc(32);  
    free(malloc(1));  
    free(a);  
}
```

```
stuXXX@spN ~/linklab/part2 $ make run test1  
[0001] Memory tracer started.  
[0002]          (nil) : malloc( 1024 ) = 0x2415060  
[0003]          (nil) : malloc( 32 ) = 0x24154c0  
[0004]          (nil) : malloc( 1 ) = 0x2415540  
[0005]          (nil) : free( 0x2415540 )  
[0006]          (nil) : free( 0x24154c0 )  
[0007]  
[0008] Statistics  
[0009]   allocated_total      1057  
[0010]   allocated_avg       352  
[0011]   freed_total         33  
[0012]  
[0013] Non-deallocated memory blocks  
[0014]   block              size      ref cnt   caller  
[0015]   0x2415060           1024        1      ??? : 0  
[0016]  
[0017] Memory tracer stopped.  
stuXXX@spN ~/linklab/part2 $
```

(Part 3) Pinpointing call locations

```
//  
// return the PC of the callsite to the dynamic memory management function  
//  
//  fname      pointer to character array to hold function name  
//  fnlen      length of character array  
//  ofs        pointer to offset to hold PC offset into function  
//  
// returns  
//  0          on success  
//  <0        on error  
//  
int get_callinfo(char *fname, size_t fnlen, unsigned long long *ofs);  
  
                                ~/linklab/part3/callinfo.h
```

```
int get_callinfo(char *fname, size_t fnlen,  
                unsigned long long *ofs)  
{  
    return -1;  
}  
  
                                ~/linklab/part3/callinfo.c
```

(Part 3) Pinpointing call locations

```
stuXXX@spN ~/linklab/part3 $ make run test1
[0001] Memory tracer started.
[0002]      main:6  : malloc( 1024 ) = 0x14f0060
[0003]      main:10 : malloc( 32 ) = 0x14f04c0
[0004]      main:1d : malloc( 1 ) = 0x14f0540
[0005]      main:25 : free( 0x14f0540 )
[0006]      main:2d : free( 0x14f04c0 )
[0007]
[0008] Statistics
[0009]   allocated_total      1057
[0010]   allocated_avg        352
[0011]   freed_total          33
[0012]
[0013] Non-deallocated memory blocks
[0014]   block                size      ref cnt    caller
[0015]   0x14f0060            1024        1      main:6
[0016]
[0017] Memory tracer stopped.
stuXXX@spN ~/linklab/part3 $
```

(Part 3) Backtracing library

- `#include <libunwind.h>`
- `int unw_get_proc_name(unw_cursor_t *cp, char *bufp, size_t len, unw_word_t *offp);`
 - Get name of current procedure
 - Returns the name of the procedure that created the stack frame identified by argument `cp`
- `int unw_getcontext(unw_context_t *ucp);`
- `int unw_init_local(unw_cursor_t *c, unw_context_t *ctxt);`
- `int unw_step(unw_cursor_t *cp);`

(Part 3) Backtracing library

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <fcntl.h>
#include <unistd.h>
#include <errno.h>

#define UNW_LOCAL_ONLY
#include <libunwind.h>
```

```
static void print_backtrace(void)
```

```
{
    //
    // having fun with functions of unwind library =)|
    //
}
```

```
void dummy(void)
```

```
{
    print_backtrace();
}
```

```
int main(void)
```

```
{
    dummy();

    return 0;
}
```

```
root@sp3:/home/ta/hkim# ./test
ip = 0x55c7345cdc82 (dummy), sp = 0x7ffca4c637e0 off=0x4
ip = 0x55c7345cdc8e (main), sp = 0x7ffca4c637f0 off=0x4
ip = 0x7f8c3c09fb97 (__libc_start_main), sp = 0x7ffca4c63800 off=0xe2
ip = 0x55c7345cdala (_start), sp = 0x7ffca4c638c0 off=0x25
```

output

objdump -d test

```
00000000000000c79 <dummy>:
c79: 55                push    %rbp
c7a: 48 89 e5          mov     %rsp,%rbp
c7d: e8 78 fe ff ff    callq   afa <print_backtrace>
c82: 90                nop
c83: 5d                pop     %rbp
c84: c3                retq

00000000000000c85 <main>:
c85: 55                push    %rbp
c86: 48 89 e5          mov     %rsp,%rbp
c89: e8 eb ff ff ff    callq   c79 <dummy>
c8e: b8 00 00 00 00    mov     $0x0,%eax
c93: 5d                pop     %rbp
c94: c3                retq
c95: 66 2e 0f 1f 84 00 00 nopw    %cs:0x0(%rax,%rax,1)
c9c: 00 00 00
c9f: 90                nop
```



(Bonus) Detect and ignore illegal deallocations

- Detect double- free / illegal free

test4.c - test case for bonus part

```
#include <stdlib.h>

void main(void) {
    void *a;

    a = malloc(1024);
    free(a);
    free(a);
    free((void*)0x1706e90);
}
```

output

```
stuXXX@spN ~/linklab/bonus $ make run test4
[0001] Memory tracer started.
[0002]      main:6   : malloc( 1024 ) = 0x1b30060
[0003]      main:11  : free( 0x1b30060 )
[0004]      main:19  : free( 0x1b30060 )
[0005]      *** DOUBLE_FREE *** (ignoring)
[0006]      main:23  : free( 0x1706e90 )
[0007]      *** ILLEGAL_FREE *** (ignoring)
[0008]
[0009] Statistics
[0010]   allocated_total      1024
[0011]   allocated_avg       1024
[0012]   freed_total         1024
[0013]
[0014] Memory tracer stopped.
stuXXX@spN ~/linklab/bonus $
```

Skeleton code snippet

```
//
// init - this function is called once when the shared library is loaded
//
__attribute__((constructor))
void init(void)
{
    char *error;

    LOG_START();

    // initialize a new list to keep track of all memory (de-)allocations
    // (not needed for part 1)
    list = new_list();

    // ...
}

//
// fini - this function is called once when the shared library is unloaded
//
__attribute__((destructor))
void fini(void)
{
    // ...

    LOG_STATISTICS(OL, OL, OL);

    LOG_STOP();

    // free list (not needed for part 1)
    free_list(list);
}
```

Utilities

- Read someone else's code
- memlog.c
 - `int mlog(int pc, const char* fmt, ...)`
- memlist.c
 - `item *new_list(void)`
 - `void free_list(void)`
 - `item *alloc(item *list, void *ptr, size_t size)`
 - `item *dealloc(item *list, void *ptr)`
 - `item *find(item *list, void *ptr)`
 - `void dump_list(item *list)`